

Class 7: machine learning for bioinformatics (part 1)

Torrey Rhyne (A14397504)

k-means clustering


We will start with k-means clustering, one of the most prevalent of all clustering methods. A disadvantage is that you need to know how many clusters there should be ($k = ?$).

Generate random data:

```
tmp <- c(rnorm(30, 3), rnorm(30, -3))  
data <- cbind(x=tmp, y = rev(tmp))  
head(data)
```

	x	y
[1,]	1.464168	-3.536458
[2,]	2.813973	-2.759121
[3,]	4.855887	-3.084354
[4,]	4.006086	-3.131239
[5,]	5.253593	-2.038764
[6,]	3.462666	-2.634887

```
plot(data)
```



The main function in R for k-means clustering is called `kmeans()`:

```
k <- kmeans(data, centers = 2, nstart = 20)
k # output is a list of "components"
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-3.243116	3.071380
2	3.071380	-3.243116

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 66.84442 66.84442
(between_SS / total_SS = 89.9 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

How many points are in each cluster?

```
k[7]
```

```
$size
```


```
[1] 30 30
```

```
k$size
```

```
[1] 30 30
```


Make a plot of our data colored by clustering results with the cluster centers shown:

```
plot(data, col = k$cluster, pch = 16)
points(k$centers, col = "blue", pch = 15, cex = 2)
```



What if we did 3 clusters?

```
k <- kmeans(data, centers = 3, nstart = 20)
plot(data, col = k$cluster, pch = 16)
points(k$centers, col = "blue", pch = 15, cex = 2)
```



Hierarchical Clustering

Hierarchical clustering has the advantage that it can reveal the structure in your data rather than imposing a structure (k-means).

The main function in base R is `hclust()`. It requires a distance matrix as its input, not the data itself:

```
hc <- hclust(dist(data))
hc # output isn't very helpful
```

Call:


```
hclust(d = dist(data))
```

```
Cluster method : complete
Distance       : euclidean
```

Number of objects: 60

Plot:

```
plot(hc)
# cut the tree
abline(h=8, col="red")
```




The function to get our clusters/groups from a hclust object is called `cutree()`:

```
groups <- cutree(hc,h=8)
```

Plot hclust results in terms of our data colored by cluster membership:

```
plot(data, col = groups)
```



Produced the same output, but we didn't need to know groups = 2 beforehand.

Principal Component Analysis (PCA)

(1) PCA of UK food data

Import data:

```
# load
url <- "https://tinyurl.com/UK-foods"
data <- read.csv(url, row.names = 1) # this is better than rownames(x) <- x[,1] approach

# examine data
str(data)
```

```
'data.frame':  17 obs. of  4 variables:
 $ England : int  105 245 685 147 193 156 720 253 488 198 ...
 $ Wales   : int  103 227 803 160 235 175 874 265 570 203 ...
 $ Scotland: int  103 242 750 122 184 147 566 171 418 220 ...
 $ N.Ireland: int  66 267 586 93 209 139 1033 143 355 187 ...
```

```
head(data)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
nrow(data)
```


```
[1] 17
```

```
ncol(data)
```


```
[1] 4
```

Try to visualize 17 variables with a bar plot:

```
barplot(as.matrix(data), beside=T, col=rainbow(nrow(data)))
```




```
barplot(as.matrix(data), beside=F, col=rainbow(nrow(data)))
```



Try generating all pairwise plots:

```
pairs(data, col=rainbow(10), pch=16)
```



```
# point on the diagonal means those 2 countries have a similar value for that variable
# the colored points represent the 17 variables (ex: "cheese")
```

These plots are difficult to make sense of. PCA to the rescue!

The main function for PCA in base R is called `prcomp()`:

```
pca <- prcomp(t(data)) # need to transpose because PCA expects variables to be columns
summary(pca)
```

Importance of components:


	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

PCA output (like k-means) is a list. One of the main PCA results people use is a 2D “score plot” (PC1 vs. PC2):

```
pca$x
```


	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
# plot just PC1 vs. PC2 (represent 96.5% of the variance)
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500), ylim=c(-300,300))
text(pca$x[,1], pca$x[,2], colnames(data))
```



Add color:

```
color <- c("orange", "red", "blue", "green")
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500), ylim=c(-300,300))
text(pca$x[,1], pca$x[,2], colnames(data), col = color)
```



Variable loadings:


Calculate how much variation in the original data each PC accounts for:

```
s <- summary(pca)
s$importance
```

	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	2.921348e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

This information can be summarized in a plot of the variances (eigenvalues) with respect to the principal component number (eigenvector number). Plot:


```
barplot(s$importance[2,], xlab="Principal Component", ylab="Percent Variation")
```



We can also consider the influence of each of the original variables upon the principal components (typically known as loading scores). This information can be obtained from the `prcomp()` returned `$rotation` component.

Plot (focus on PC1 since it accounts for >90% of the variance):


```
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



Here we see variables (foods) with the largest positive / negative “loading scores” that effectively “push” observations (countries) to right / left side of the plot.

Generate a similar ‘loadings plot’ for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```




The largest positive loading score = soft drinks, pushing Scotland to the top of the plot. The largest negative loading score = fresh potatoes, pushing Wales to the bottom of the plot.


Let's make the score plot figure look nicer with ggplot:

```
library(ggplot2)

df <- as.data.frame(pca$x)
df_lab <- tibble::rownames_to_column(df, "Country")
# first basic plot
ggplot(df_lab) +
  aes(PC1, PC2, col=Country) +
  geom_point()
```




```
# nicer plot
ggplot(df_lab) +
  aes(PC1, PC2, col=Country, label=Country) +
  geom_hline(yintercept = 0, col="gray") +
  geom_vline(xintercept = 0, col="gray") +
  geom_point(show.legend = FALSE) +
  geom_label(hjust=1, nudge_x = -10, show.legend = FALSE) +
  expand_limits(x = c(-300,500)) +
  xlab("PC1 (67.4%)") +
  ylab("PC2 (28%)") +
  theme_bw()
```




Let's do the same for our loadings/PC contributions figures:

```
ld <- as.data.frame(pca$rotation)
ld_lab <- tibble::rownames_to_column(ld, "Food")

# first basic plot
ggplot(ld_lab) +
  aes(PC1, Food) +
  geom_col()
```

```
# nicer plot
ggplot(ld_lab) +
  aes(PC1, reorder(Food, PC1), bg=PC1) +
  geom_col() +
  xlab("PC1 Loadings/Contributions") +
  ylab("Food Group") +
  scale_fill_gradient2(low="purple", mid="gray", high="darkgreen", guide=NULL) +
  theme_bw()
```



Another way to see this information together with the main PCA plot is in a so-called biplot:

```
## can be useful for small datasets  
biplot(pca)
```

